

# Hybrid Circuit/Packet Network Scheduling with Multiple Composite Paths

---

Shih-Hao Tseng<sup>1</sup>, (pronounced as “She-How Zen”)  
joint work with Bo Bai<sup>2</sup> and John C. S. Lui<sup>3</sup>

April 18, 2018


<sup>1</sup>School of Electrical and Computer Engineering, Cornell University

<sup>2</sup>Future Network Theory Lab, 2012 Labs, Huawei Technologies, Co. Ltd.

<sup>3</sup>Department of Computer Science and Engineering, The Chinese University of Hong Kong


# Hybrid Switches

- A hybrid switch (h-switch) combines an electronic packet switch (EPS) and an optical circuit switch (OCS).




# Hybrid Switches

- A hybrid switch (h-switch) combines an electronic packet switch (EPS) and an optical circuit switch (OCS).
- EPS can switch among many-to-many routing patterns swiftly; and OCS provides high-bandwidth one-to-one routing.




# Drawbacks of Hybrid Switches

- EPS suffers low bandwidth;




# Drawbacks of Hybrid Switches

- EPS suffers low bandwidth; and OCS suffers slow time division multiplexing (TDM) when mapping many-to-one or one-to-many.




# Drawbacks of Hybrid Switches

- EPS suffers low bandwidth; and OCS suffers slow time division multiplexing (TDM) when mapping many-to-one or one-to-many.
- The drawbacks restrict the use of hybrid switches in data-parallel applications, such as MapReduce.




# Composite Paths

- Since EPS nowadays supports heterogeneous port bandwidth, with many low-bandwidth ports and few high bandwidth ports, one can connect an OCS output to an EPS input (and vice versa) to create a *composite path* (Vargaftik et al., 2016).



# Advantages of Composite Paths


- Composite paths allow EPS to send more data to the outputs under many-to-one mapping.






# Advantages of Composite Paths

- Composite paths allow EPS to send more data to the outputs under many-to-one mapping.
- Composite paths avoid OCS TDM but still provide higher input bandwidth for one-to-many scenarios.




# Challenges of Composite-Path Scheduling

- Without composite paths, EPS and OCS can be scheduled in parallel (h-switch).




# Challenges of Composite-Path Scheduling

- Without composite paths, EPS and OCS can be scheduled in parallel (h-switch).
- However, with composite paths, EPS and OCS are tangled together (cp-switch).




# Demand Reduction

- Vargaftik et al. (2016) suggested some heuristics to translate the demand matrix so that we can schedule cp-switches using h-switch scheduling algorithms.




# Unresolved Issues

- The translation based algorithm does not provide a theoretical performance guarantee.



# Unresolved Issues


- The translation based algorithm does not provide a theoretical performance guarantee.
- The algorithm only works for one pair of composite paths. In general, we can have more composite paths at the system level.



# Composite-Path Switch Scheduling


Goal: finding a shortest schedule to satisfy i/o demand.

- Systematic analysis of cp-switch schedules.
- Performance guarantee for cp-switch scheduling algorithms.
- Applicability to multiple composite paths.



# Formulation


- $N$  ports. Each port  $n \in [1, N]_{\mathbb{Z}} = \{1, \dots, N\}$  connects to both EPS and OCS.
- $P$  composite paths (each composite path is a full-duplex line connecting EPS and OCS).






# Formulation

- Each port is assumed to have symmetric input/output capacity (bandwidth).
- Each EPS port has capacity  $c_E$ ; Each OCS port has capacity  $c_O$ ; And each composite path  $p \in P$  is assumed to have capacity  $c_O$  as well. In general,  $c_O \sim 10c_E \gg c_E$ .




# Formulation

- Each OCS input maps to at most one OCS output, and each OCS output can receive data from at most one OCS input. Such mapping is called an *OCS configuration*.
- No data can be buffered at the inputs or the outputs of the composite paths.




# Formulation

- Each entry  $D_{ij}$  in the demand matrix  $D \in \mathbb{R}^{N \times N}$  refers to the amount of data that should be sent from port  $i$  to port  $j$ .
- By convention, we assume  $D_{nn} = 0$  for each  $n \in [1, N]_{\mathbb{Z}}$ .




# Formulation

- The reconfiguration time of OCS is  $\delta$ . During the reconfiguration, OCS stops carrying data. In contrast, EPS changes the sending rate seamlessly.




# Formulation

- A  $M + 1$  step schedule is considered: In step 0, only EPS is used; The remaining  $M$  steps involve the whole cp-switch.
- Each step  $m \in [1, M]_{\mathbb{Z}}$  consists of a reconfiguration phase and a sending phase. The length of the step  $m$  is  $\tau^m$ .




# Continuous-Time Control Formulation

- Let  $\mathcal{E}(t), \mathcal{U}(t), \mathcal{V}(t), \mathcal{O}(t) \in \mathbb{R}^{N \times N}$  be the *mapping matrices*, which represent the port-to-port sending rates, at time  $t$ .
- We can formulate the scheduling problem as a continuous-time control problem.



# Mixed Integer Linear Programming (MILP) Formulation

- It turns out the continuous-time control problem can be equivalently transformed into an MILP.
- Instead of the continuous-time mapping matrices, we express the problem in terms of the total data sent during each phase.



# Mixed Integer Linear Programming (MILP) Formulation

- For each  $\hat{M} \in [0, M]_{\mathbb{Z}}$ , solving the following subproblems:

$$I(\hat{M}) = \min \text{ length of } (\hat{M} + 1)\text{-step schedule}$$

s.t. parameter setup

demand constraints

capacity constraints

operation constraints

leads to the shortest schedule  $\text{OPT} = \min_{\hat{M} \in [0, M]_{\mathbb{Z}}} I(\hat{M})$ .



# Mixed Integer Linear Programming (MILP) Formulation

- For each  $\hat{M} \in [0, M]_{\mathbb{Z}}$ , solving the following subproblems:

$$I(\hat{M}) = \min \sum_{m=0}^{\hat{M}} \tau^m$$

s.t. parameter setup

demand constraints

capacity constraints

operation constraints

leads to the shortest schedule  $\text{OPT} = \min_{\hat{M} \in [0, M]_{\mathbb{Z}}} I(\hat{M})$ .

# Mixed Integer Linear Programming (MILP) Formulation

- For each  $\hat{M} \in [0, M]_{\mathbb{Z}}$ , solving the following subproblems:

$$I(\hat{M}) = \min \sum_{m=0}^{\hat{M}} \tau^m$$

s.t. parameter setup

$$E^0 + \sum_{m=1}^{\hat{M}} E^m + U^m + V^m + O^m = D$$

capacity constraints

operation constraints

leads to the shortest schedule  $\text{OPT} = \min_{\hat{M} \in [0, M]_{\mathbb{Z}}} I(\hat{M})$ .

# Mixed Integer Linear Programming (MILP) Formulation

- For each  $\hat{M} \in [0, M]_{\mathbb{Z}}$ , solving the following subproblems:

$$I(\hat{M}) = \min \sum_{m=0}^{\hat{M}} \tau^m$$

s.t. parameter setup

$$E^0 + \sum_{m=1}^{\hat{M}} E^m + U^m + V^m + O^m = D$$

capacity constraints


OCS must map one-to-one

leads to the shortest schedule  $\text{OPT} = \min_{\hat{M} \in [0, M]_{\mathbb{Z}}} I(\hat{M})$ .

# Composite-Path Switch Scheduling

Goal: finding a shortest schedule to satisfy i/o demand.

- ✓ Systematic analysis of cp-switch schedules.
- Performance guarantee for cp-switch scheduling algorithms.
- Applicability to multiple composite paths.



- Unfortunately,  $I(\hat{M})$  is NP-hard.

$$I(\hat{M}) = \min \sum_{m=0}^{\hat{M}} \tau^m$$

s.t. parameter setup

$$E^0 + \sum_{m=1}^{\hat{M}} E^m + U^m + V^m + O^m = D$$

capacity constraints

OCS must map one-to-one

 NP-hardness

# Linear Relaxation

- We then linear-relax  $I(\hat{M})$  to be  $L(\hat{M})$ :

$$L(\hat{M}) = \min \sum_{m=0}^{\hat{M}} \tau^m$$

s.t. parameter setup

$$E^0 + \sum_{m=1}^{\hat{M}} E^m + U^m + V^m + O^m = D$$

capacity constraints

OCS can map many-to-many

## Linear Relaxation: Special Case

- The schedule that uses EPS only :

$$L(0) = \min \tau^0$$

s.t. parameter setup

$$E^0 = D$$

capacity constraints

- Especially,  $L(0) = I(0)$ .

## Lemma 1

*Any cp-switch scheduling algorithm adopting  $L(0)$  as an upper bound is a  $\frac{c_E + c_O}{c_E}$ -approximation algorithm.*


- Lemma 1 implies that a scheduling algorithm that produces shorter schedule than EPS only schedule is an approximation algorithm with approximation ratio  $\frac{c_E + c_O}{c_E}$ .  
 $\Rightarrow$  Comparing with  $L(0)$  is a naive way to have performance guarantee.



# Composite-Path Switch Scheduling

Goal: finding a shortest schedule to satisfy i/o demand.

- ✓ Systematic analysis of cp-switch schedules.
- ✓ Performance guarantee for cp-switch scheduling algorithms.
- Applicability to multiple composite paths.



## Lemma 2


*If  $L(0) \leq \delta$ ,  $\text{OPT} = L(0)$  and the shortest time schedule uses EPS only. Otherwise,*

$$L(0) \geq \text{OPT} \geq L(1).$$

- Lemma 2 inspires us to find the shortest time schedule “between”  $L(1)$  and  $L(0)$ .  
 $\Rightarrow$  Since  $L(1)$  is not feasible to  $I(\hat{M})$ , can we round it to a feasible schedule shorter than  $L(0)$ ?

# Uprounding Procedure


- Taking the demand  $D$ , we compute two schedules  $L(0)$  and  $L(1)$ .




Demand matrix  $D$

# Uprounding Procedure

- Taking the demand  $D$ , we compute two schedules  $L(0)$  and  $L(1)$ .




1-step linear-relaxed schedule  $L(0)$



2-step linear-relaxed schedule  $L(1)$

# Uprounding Procedure

- Taking the demand  $D$ , we compute two schedules  $L(0)$  and  $L(1)$ .
- If  $L(0) \leq \delta$ , we have found the shortest schedule.





1-step linear-relaxed schedule  $L(0)$



# Uprounding Procedure


- Taking the demand  $D$ , we compute two schedules  $L(0)$  and  $L(1)$ .
- If  $L(0) \leq \delta$ , we have found the shortest schedule.
- Otherwise, we upround  $L(1)$  to a feasible schedule.



2-step linear-relaxed schedule  $L(1)$

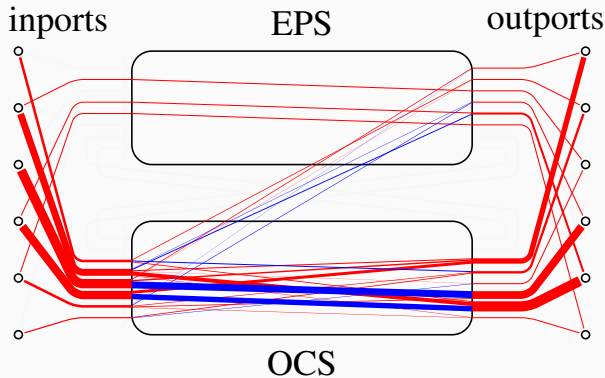
# Uprounding Procedure

- Find OCS configuration that can send as much relaxed traffic as possible.



# Uprounding Procedure


- Find OCS configuration that can send as much relaxed traffic as possible.  
⇒ Use maximum weight matching algorithm.






# Uproounding Procedure

- The OCS Configuration that supports the most relaxed traffic.




# Uprounding Procedure

- The OCS Configuration that supports the most relaxed traffic.




## Uprounding Procedure

- Once the OCS configuration is decided, the shortest 2-step schedule can be obtained by a linear program.




# Uprounding Procedure

- Once the OCS configuration is decided, the shortest 2-step schedule can be obtained by a linear program.





# Uprounding Procedure

- Once the OCS configuration is decided, the shortest 2-step schedule can be obtained by a linear program.





# Uprounding Procedure

- Once the OCS configuration is decided, the shortest 2-step schedule can be obtained by a linear program.




# Uprounding Procedure

- Once the OCS configuration is decided, the shortest 2-step schedule can be obtained by a linear program.  
⇒ The algorithm can work online by repeating the same procedure on the resulting step 0.



# Bound Shrinking

- Let  $\tau_R$  be the current length of the schedule without step 0.
- Upper bound:  $L(0) + \tau_R$ .
- Lower bound:  $L(1) + \tau_R$ .






# Composite-Path Switch Scheduling

Goal: finding a shortest schedule to satisfy i/o demand.

- ✓ Systematic analysis of cp-switch schedules.
- ✓ Performance guarantee for cp-switch scheduling algorithms.
- ✓ Applicability to multiple composite paths.




## Issues to be Evaluated


- Do more composite paths lead to shorter schedule?
- How well do the algorithms schedule one-to-many traffic?
- How does OCS reconfiguration time  $\delta$  influence the length of the schedule?

# Benefits of Multiple Composite Paths

- Meshed multicast traffic: Pick each input-output pair with probability 0.5.
- Skewed multicast traffic: Pick each source with probability 0.5 to install multicast traffic.




Meshed multicast traffic.




Skewed multicast traffic.

# Benefits of Multiple Composite Paths



(a) Meshed multicast traffic.




(b) Skewed multicast traffic.

**Figure 1:** The 30<sup>th</sup>, 40<sup>th</sup>, 50<sup>th</sup>, 60<sup>th</sup>, and 70<sup>th</sup> percentiles of the schedule length given by the proposed algorithm.


# Performance Improvement under Skewed Demand

- We compare our algorithm with the state-of-the-art scheduling algorithm CPSwitchSched (Vargaftik et al., 2016).




# Performance Improvement under Skewed Demand

- We compare our algorithm with the state-of-the-art scheduling algorithm CPSwitchSched (Vargaftik et al., 2016).





Traffic through EPS and OCS



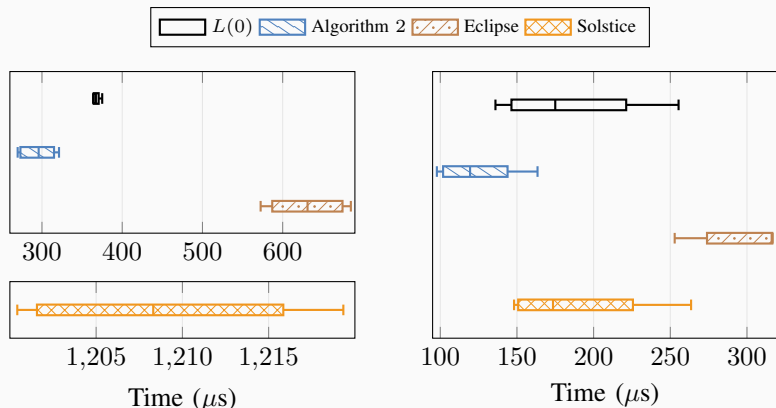
Traffic through the composite path

# Performance Improvement under Skewed Demand

- We compare our algorithm with the state-of-the-art scheduling algorithm CPSwitchSched (Vargaftik et al., 2016).
  - Solstice (Liu et al., 2015).
  - Eclipse (Venkatakrisnan et al., 2016).



# Performance Improvement under Skewed Demand



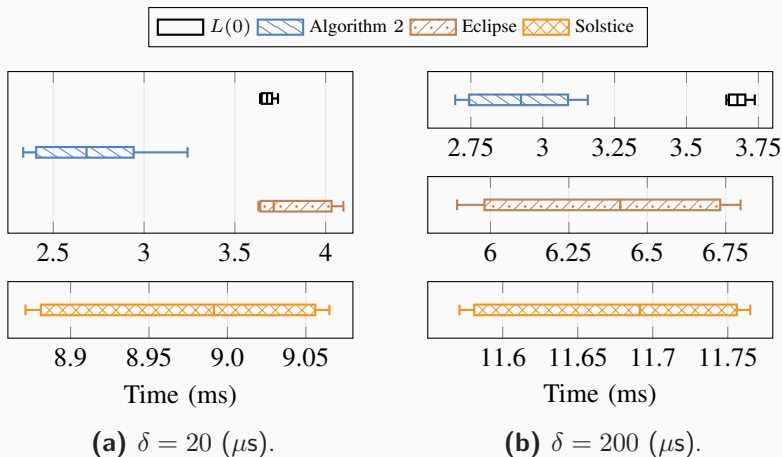
(a) Meshed traffic,  $\delta = 20 \mu s$ .

(b) Skewed traffic,  $\delta = 20 \mu s$ .

**Figure 4:** The 1<sup>st</sup>-5<sup>th</sup>-50<sup>th</sup>-95<sup>th</sup>-99<sup>th</sup> percentiles of the schedule lengths.

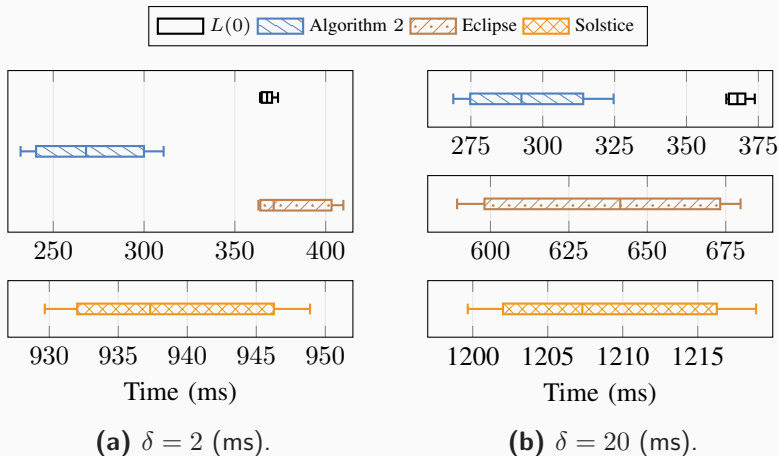


# Effects of OCS Reconfiguration Overhead



**Figure 5:** Lighter Loading Condition.

# Effects of OCS Reconfiguration Overhead



**Figure 6:** Heavier Loading Condition.

# Effects of OCS Reconfiguration Overhead

Loading, $\delta$	Solstice	Eclipse
Lighter, 20 ( $\mu$ s)	70.2%	27.2%
Lighter, 200 ( $\mu$ s)	75.0%	54.4%
Heavier, 2 (ms)	71.4%	27.8%
Heavier, 20 (ms)	75.8%	54.4%



**Table 1:** Performance improvement of the given algorithm on the 50<sup>th</sup> percentile over CPSwitchSched with different schedulers.

# Conclusion


- We establish a framework to study cp-switch scheduling problems systematically. It supports multiple composite paths.
- Although each cp-switch scheduling subproblem is NP-hard, a fixed approximation ratio is still possible for the overall schedule.
- Our proposed algorithm not only works online but also outperforms existing methods significantly (by 30% to 70%).

## Questions & Answers

## References

-  S. Vargaftik, K. Barabash, Y. Ben-Itzhak, O. Biran, I. Keslassy, D. Lorenz, and A. Orda, “Composite-path switching,” in *Proc. CoNEXT*. ACM, 2016, pp. 329–343.
-  H. Liu, M. K. Mukerjee, C. Li, N. Feltman, G. Papen, S. Savage, S. Seshan, G. M. Voelker, D. G. Anderson, M. Kaminsky, G. Porter, and A. C. Snoeren, “Scheduling techniques for hybrid circuit/packet networks,” in *Proc. CoNEXT*. ACM, 2015, pp. 41:1–41:13.

# References

-  S. Bojja Venkatakrisnan, M. Alizadeh, and P. Viswanath,  
“Costly circuits, submodular schedules and approximate  
carathéodory theorems,” in *ACM SIGMETRICS*, vol. 44, no. 1.  
ACM, 2016, pp. 75–88.